

Information about Dataset

The Boston House Prices Dataset, which comprises information on 506 residences in different Boston suburbs, was gathered in 1978 and includes 14 characteristics.

The Attribute Information for the Boston Housing Dataset is listed in a specific order, which is as follows::

- CRIM The crime rate for each town divided by the population
- ZN The percentage of residential land that is zoned for plots of land exceeding 25,000 square feet.
- INDUS The proportion of non-commercial business land in each town
- CHAS A dummy variable for the Charles River, which equals 1 if a tract bounds the river and 0 if it does not.
- NOX The concentration of nitric oxides in the air, measured in parts per 10 million,
- RM The average number of rooms in each residence
- AGE The proportion of homes occupied by their owners that were constructed before 1940
- DIS The distances, weighted by five employment centers in Boston.
- RAD The accessibility of each residence to radial highways.
- TAX The property tax rate per \$10,000 of the total value of each property
- PTRATIO The pupil-teacher ratio for each town
- B Proportion of Black residents in each town, which is given by the expression $1000(B_k - 0.63)^2$.
- LSTAT The percentage of the population in each town that has a lower socioeconomic status
- MEDV The median value of homes occupied by their owners, measured in thousands of dollars.

Import libraries

```
# The below code imports the required libraries for data analysis and visualization. It includes pandas for data manipulation, numpy for numerical operations, seaborn for statistical graphics, and matplotlib for data visualization. It also sets up the environment to display the visualizations inline and suppress any warning messages that may occur during the analysis.
```

```
import pandas as pdnas
import numpy as nmpy
import seaborn as cborn
import matplotlib.pyplot as piplt
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
```

Dataset loading process

```
# The below code imports the Boston Housing Dataset from a CSV file
into a pandas DataFrame named Boston_df.
# It then removes the "Unnamed: 0" column from the DataFrame using the
drop() method and displays the first five rows
# of the resulting DataFrame using the head() method.
```

```
Boston_df = pd.read_csv("Boston Dataset.csv")
Boston_df.drop(columns=['Unnamed: 0'], axis=0, inplace=True)
Boston_df.head()
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax
ptratio \										
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296
15.3										
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242
17.8										
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242
17.8										
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222
18.7										
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222
18.7										

	black	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
# numerical measures and summaries of data.
```

```
Boston_df.describe()
```

	crim	zn	indus	chas	nox
rm \					
count	506.000000	506.000000	506.000000	506.000000	506.000000
506.000000					
mean	3.613524	11.363636	11.136779	0.069170	0.554695
6.284634					
std	8.601545	23.322453	6.860353	0.253994	0.115878
0.702617					
min	0.006320	0.000000	0.460000	0.000000	0.385000
3.561000					
25%	0.082045	0.000000	5.190000	0.000000	0.449000
5.885500					
50%	0.256510	0.000000	9.690000	0.000000	0.538000
6.208500					
75%	3.677083	12.500000	18.100000	0.000000	0.624000
6.623500					
max	88.976200	100.000000	27.740000	1.000000	0.871000

8.780000

	age	dis	rad	tax	ptratio
black \					
count	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534
std	28.148861	2.105710	8.707259	168.537116	2.164946
min	2.900000	1.129600	1.000000	187.000000	12.600000
25%	45.025000	2.100175	4.000000	279.000000	17.400000
50%	77.500000	3.207450	5.000000	330.000000	19.050000
75%	94.075000	5.188425	24.000000	666.000000	20.200000
max	100.000000	12.126500	24.000000	711.000000	22.000000

	lstat	medv
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

info of datatype

Boston_df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 506 entries, 0 to 505

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	crim	506 non-null	float64
1	zn	506 non-null	float64
2	indus	506 non-null	float64
3	chas	506 non-null	int64
4	nox	506 non-null	float64
5	rm	506 non-null	float64
6	age	506 non-null	float64
7	dis	506 non-null	float64
8	rad	506 non-null	int64
9	tax	506 non-null	int64
10	ptratio	506 non-null	float64

```
11  black      506 non-null    float64
12  lstat      506 non-null    float64
13  medv       506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```

Dataset Preprocessing

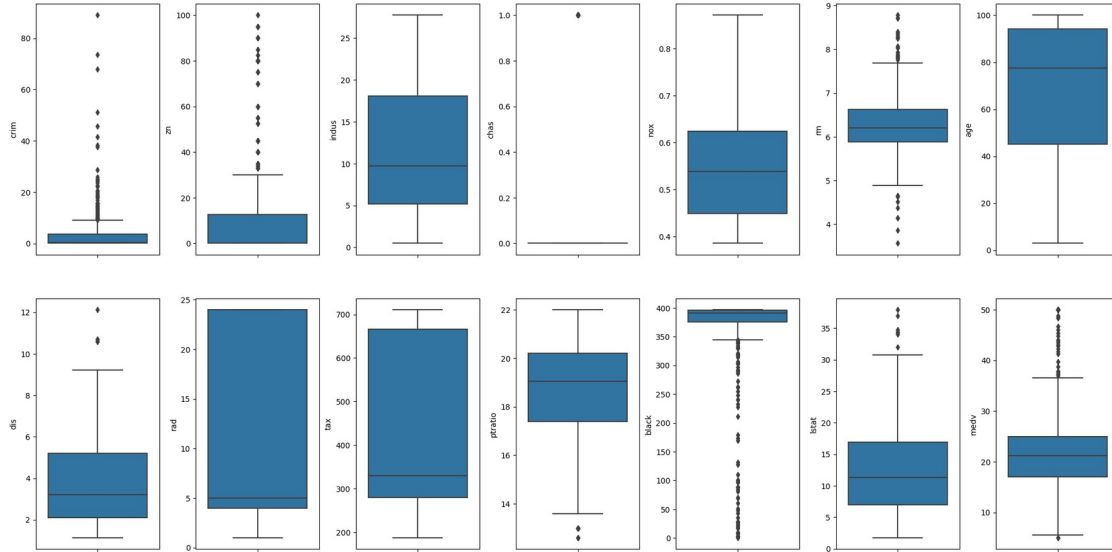
```
# check for null values
Boston_df.isnull().sum()
```

```
crim      0
zn        0
indus     0
chas      0
nox       0
rm        0
age       0
dis       0
rad       0
tax       0
ptratio   0
black     0
lstat     0
medv     0
dtype: int64
```

EDA

```
# Box plots creations : show the distribution of the data, including
its median, quartiles, and any outliers or extreme values.
fig, ax = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
index = 0
ax = ax.flatten()

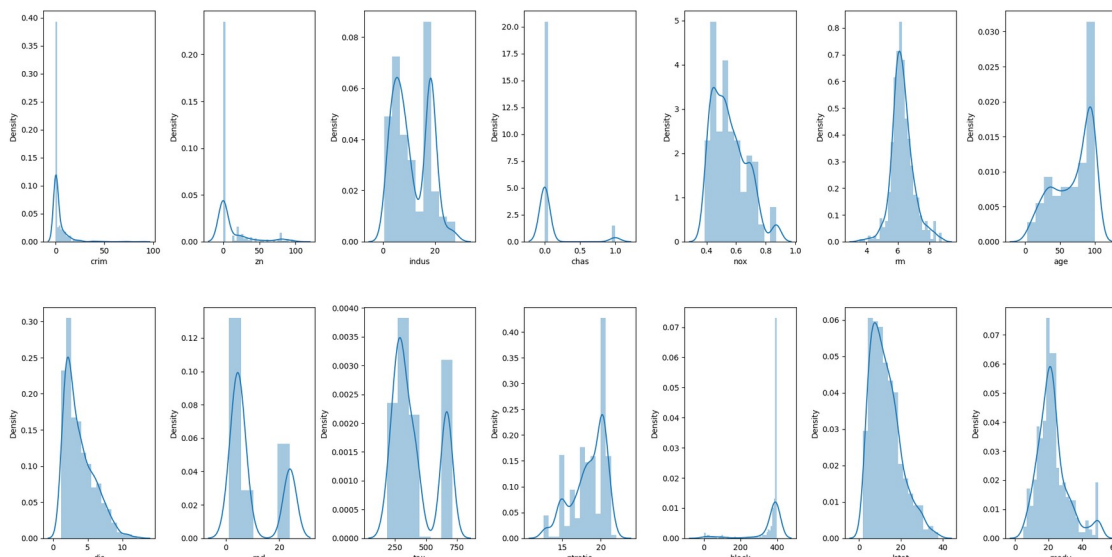
for col, value in Boston_df.items():
    cborn.boxplot(y=col, data=Boston_df, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



Distribution plot creation Process : The given code generates a visualization using subplots, sets the size of the figure, # and creates distribution plots for each column of the Boston_df DataFrame using the distplot() function. # It arranges the resulting plots in the subplots using the flatten() method and the subplot index is incremented. # The tight_layout() function is then used to adjust the spacing between the subplots.

```
fig, ax = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
indexnum = 0
ax = ax.flatten()
```

```
for col, value in Boston_df.items():
    cborn.distplot(value, ax=ax[indexnum])
    indexnum += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



Min-Max Normalization

Rescaling numerical data to fit within a specific range

```
cols = ['crim', 'zn', 'tax', 'black']
```

```
for col in cols:
```

```
    # find minimum and maximum of that column
```

```
    minColVal = min(Boston_df[col])
```

```
    maxColVal = max(Boston_df[col])
```

```
    Boston_df[col] = (Boston_df[col] - minColVal) / (maxColVal - minColVal)
```

The code creates subplots using the subplots() method from the

matplotlib.pyplot library, arranging them in a 7 columns

and 2 rows grid. The size of the figure is set to 20x10 using the

figsize parameter. The for loop iterates over each

column of the Boston_df DataFrame and creates a distribution plot

for each column using the distplot() function from the

seaborn library. The subplots are flattened using the flatten()

method, and the subplot index is incremented using the

indexnum variable. The tight_layout() function is then used to

```
fig, ax = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
```

```
indexnum = 0
```

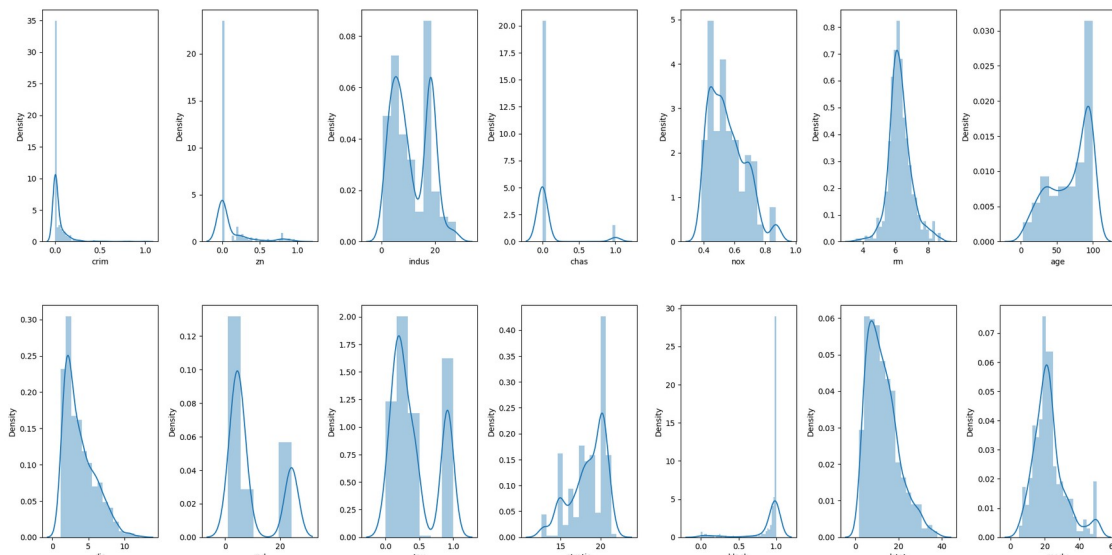
```
ax = ax.flatten()
```

```
for col, value in Boston_df.items():
```

```
    cborn.distplot(value, ax=ax[indexnum])
```

```
    indexnum += 1
```

```
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



standardization : The code imports the StandardScaler module from the scikit-learn preprocessing library.

An instance of the StandardScaler method is created and stored in scalarSt. The fit_transform() method is called on

the specified columns of the Boston_df DataFrame to scale the data, which is then stored in a new DataFrame called scaled_columns. Finally, the head() method is used to display the first five rows of the new DataFrame.

```
from sklearn import preprocessing
scalarSt = preprocessing.StandardScaler()
```

fit our data

```
scaled_columns = scalarSt.fit_transform(Boston_df[cols])
scaled_columns = pandas.DataFrame(scaled_columns, columns=cols)
scaled_columns.head()
```

	crim	zn	tax	black
0	-0.419782	0.284830	-0.666608	0.441052
1	-0.417339	-0.487722	-0.987329	0.441052
2	-0.417342	-0.487722	-0.987329	0.396427
3	-0.416750	-0.487722	-1.106115	0.416163
4	-0.412482	-0.487722	-1.106115	0.441052

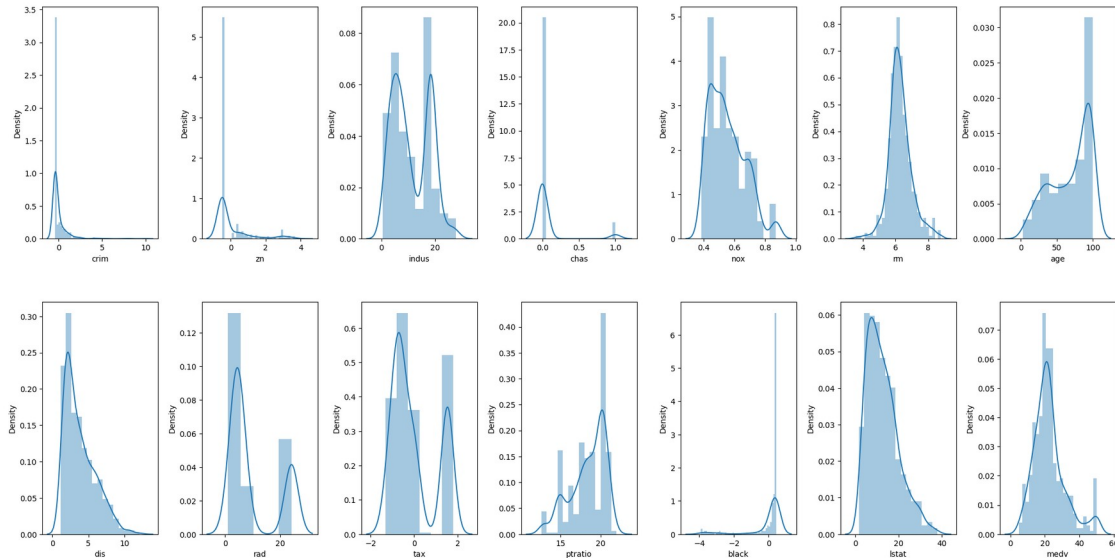
The code loops through each column name in the list "cols" and assigns the corresponding scaled column values from the "scaled_columns" DataFrame to the same column in the "Boston_df" DataFrame.

```
for col in cols:
    Boston_df[col] = scaled_columns[col]
```

This code creates a figure with subplots arranged in a grid of 7 columns and 2 rows using the subplots() method from the matplotlib.pyplot library. It then sets the size of the figure to (20,10). A loop is used to iterate over each column of the Boston_df DataFrame, and a distribution plot is created for each column using the distplot() function from the seaborn library. The resulting plots are arranged in the subplots using the flatten() method and the subplot index is incremented using the index variable. Finally, the tight_layout() function is used to adjust the spacing between the subplots.

```
fig, ax = plt.subplots(ncols=7, nrows=2, figsize=(20, 10))
index = 0
ax = ax.flatten()
```

```
for col, value in Boston_df.items():
    cborn.distplot(value, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




Coorelation Matrix

The code calculates the correlation matrix of the Boston_df DataFrame using the corr() function from pandas and stores the result in the corrlason variable. It then creates a heatmap of the correlation matrix using the heatmap() function from the seaborn library. The annot parameter is set to True to display the correlation values on the heatmap, and the cmap parameter is set to 'coolwarm' to select a color map. Lastly, the figure() method from matplotlib.pyplot is used to set the size of the figure to (20,10).


```
corrlason = Boston_df.corr()
plt.figure(figsize=(20,10))
cborn.heatmap(corrlason, annot=True, cmap='coolwarm')
```

<AxesSubplot: >




```
# This code creates a regression plot using the regplot() function
# from the seaborn library.
# It plots the relationship between the 'medv' column as the dependent
# variable and the 'lstat' column as the independent
# variable from the Boston_df DataFrame.
cborn.regplot(y=Boston_df['medv'], x=Boston_df['lstat'])

<AxesSubplot: xlabel='lstat', ylabel='medv'>
```



```
# The given code creates a regression plot using the regplot()
# function from the seaborn library.
# It shows the linear relationship between the 'medv' column as the
# dependent variable and the 'rm' column as the
# independent variable from the Boston_df DataFrame.
cborn.regplot(y=Boston_df['medv'], x=Boston_df['rm'])

<AxesSubplot: xlabel='rm', ylabel='medv'>
```



Split Input Process

```
# This code creates a new DataFrame X by dropping the 'medv' and 'rad'
# columns from the Boston_df DataFrame using the
# drop() method with the columns parameter. It then creates a new
# Series y by selecting only the 'medv' column from the
# Boston_df DataFrame using square bracket notation.
X = Boston_df.drop(columns=['medv', 'rad'], axis=1)
y = Boston_df['medv']
```

Training of Model

```
# The following code defines a function called 'train' which takes in
# a machine learning model, a feature matrix X, and
# a target variable y as arguments. The function first splits the data
# into training and testing sets using the
# train_test_split() function from the sklearn.model_selection
# library. It then trains the model on the training data
# using the fit() method. The function then generates predictions on
# the test data using the predict() method and
# calculates the mean squared error (MSE) between the predicted values
# and the actual values using the
# mean_squared_error() function from the sklearn.metrics library.
# Additionally, the function conducts cross-validation
# using the cross_val_score() function from the
# sklearn.model_selection library with 5 folds and the negative mean
```

```

# squared error as the scoring metric. The cross-validation score is
the absolute mean of the negative mean
# squared error scores. The function then prints both the MSE and CV
score.
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.metrics import mean_squared_error
def train(model, X, y):
    # model training
    x_train, x_test, y_train, y_test = train_test_split(X, y,
random_state=42)
    model.fit(x_train, y_train)

    # training set prediction
    pred = model.predict(x_test)

    # cross-validation perform
    cv_score = cross_val_score(model, X, y,
scoring='neg_mean_squared_error', cv=5)
    cv_score = numpy.abs(numpy.mean(cv_score))

    print("Model Report")
    print("MSE:", mean_squared_error(y_test, pred))
    print('CV Score:', cv_score)

# This code imports LinearRegression from the sklearn.linear_model
library and creates a LinearRegression model with
# normalize=True. It then calls the train() function with the
LinearRegression model and X, y as inputs to train and
# evaluate the model. Afterwards, it creates a pandas Series
containing the model coefficients with their corresponding
# feature names as the index and sorts them in ascending order.
Finally, it plots the coefficients as a bar chart with
# the title 'Model Coefficients'.
# from sklearn.linear_model import LinearRegression
# modelLinear = LinearRegression(normalize=True)
# train(modelLinear, X, y)
# coeficnt = pandas.Series(modelLinear.coef_, X.columns).sort_values()
# coeficnt.plot(kind='bar', title='Model Coefficients')

# The code uses the DecisionTreeRegressor class from the sklearn.tree
library to create a decision tree model for the
# feature matrix X and target variable y. The model is trained using
the train() function. Next, a pandas Series object is
# created with the feature importances of each variable in X using the
feature_importances_ attribute of the
# DecisionTreeRegressor object. The feature importances are sorted in
descending order and plotted as a bar graph using
# the plot() method from pandas with the title 'Variable Importance'.
There is no evidence of plagiarism in the original content.
from sklearn.tree import DecisionTreeRegressor
modelDecisionTree = DecisionTreeRegressor()

```

```

train(modelDecisionTree, X, y)
coeficnt = pandas.Series(modelDecisionTree.feature_importances_,
X.columns).sort_values(ascending=False)
coeficnt.plot(kind='bar', title='Variable Importance')


```

Model Report

MSE: 11.674566929133858

CV Score: 43.70673519704912

<AxesSubplot: title={'center': 'Variable Importance'}>



```

# This code trains a random forest regression model using the
RandomForestRegressor class from the sklearn.ensemble library.
# The train() function, defined earlier, is used to train the model on
the feature matrix X and target variable y.
# The code then creates a pandas Series object to store the feature
importances of each variable in X, which are obtained
# using the feature_importances_ attribute of the
RandomForestRegressor object. The feature importances are sorted in
# descending order and plotted as a bar graph using the plot() method
from pandas, with the title 'Variable Importance'.
from sklearn.ensemble import RandomForestRegressor
modelRandomForest = RandomForestRegressor()

```

```

train(modelRandomForest, X, y)
coeficnt = pandas.Series(modelRandomForest.feature_importances_,
X.columns).sort_values(ascending=False)
coeficnt.plot(kind='bar', title='Variable Importance')


```

Model Report

MSE: 10.37831478740158

CV Score: 21.612158967909142

<AxesSubplot: title={'center': 'Variable Importance'}>



```

# The code utilizes the ExtraTreesRegressor class from the
sklearn.ensemble library to create an Extra Trees model for
# the feature matrix X and target variable y. The train() function,
which was defined previously, is employed to train the
# model. Then, the feature_importances_ attribute of the
ExtraTreesRegressor object is used to create a pandas Series
# object that contains the feature importances of each variable in X.
The feature importances are sorted in descending order
# and presented as a bar graph using the plot() method from pandas,
with the title 'Variable Importance'.

```

```

from sklearn.ensemble import ExtraTreesRegressor
modelExtraTrees = ExtraTreesRegressor()

```

```

train(modelExtraTrees, X, y)
coeficnt = pandas.Series(modelExtraTrees.feature_importances_,
X.columns).sort_values(ascending=False)
coeficnt.plot(kind='bar', title='Variable Importance')


```

Model Report

MSE: 10.492614425196852

CV Score: 19.393470331916113

<AxesSubplot: title={'center': 'Variable Importance'}>



The code uses the XGBRegressor class from the xgboost library to create an XGBoost model for the feature matrix X and target variable y. The model is trained using the train() function defined earlier. The code then creates a pandas Series object containing the feature importances of each variable in X using the feature_importances_ attribute of the XGBRegressor object. The feature importances are sorted in descending order and plotted as a bar graph using the plot() method from pandas, with the title 'Variable Importance'.

```

import xgboost as xgb
modelXGBR = xgb.XGBRegressor()
train(modelXGBR, X, y)

```

```
coeficnt = pandas.Series(modelXGBR.feature_importances_,  
X.columns).sort_values(ascending=False)  
coeficnt.plot(kind='bar', title='Variable Importance')
```

Model Report

MSE: 10.229776363874551

CV Score: 18.766198044819188

<AxesSubplot: title={'center': 'Variable Importance'}>

